


GRADIENT DESCENT


Overview



- in many algorithms we want to compute weights w to minimize cost function $J(w)$

figure reprinted from www.kdnuggets.com with explicit permission of the editor


Basic Idea



- in gradient descent, we start with initial weights and update weights to reduce $J(w)$

figure reprinted from www.kdnuggets.com with explicit permission of the editor


How to Update Weights?



- take steps proportional to the gradient

figure reprinted from www.kdnuggets.com with explicit permission of the editor


Why Use Gradient?



- cost function $J(W)$ decreases fastest in the direction of negative gradient

figure reprinted from www.kdnuggets.com with explicit permission of the editor


Intuition



- find temperature (“weights”) to maximize “enjoyment”

figure reprinted from www.kdnuggets.com with explicit permission of the editor

Equivalent Formulation



- find temperature (“weights”) to minimize “suffering” $J(w)$

figure reprinted from www.kdnuggets.com with explicit permission of the editor


Exhaustive Search



- can examine "all" values
- this is inefficient

figure reprinted from www.kdnuggets.com with explicit permission of the editor


Gradient Descent



- iteratively update weights to lower $J(w)$

figure reprinted from www.kdnuggets.com with explicit permission of the editor


Typical Step



- $J(w)$ is lowered if we move “opposite” gradient

figure reprinted from www.kdnuggets.com with explicit permission of the editor


Typical Step (cont'd)



- continue moving "opposite" gradient

figure reprinted from www.kdnuggets.com with explicit permission of the editor


”Speed” of Convergence



- can take larger step for ”steeper” slopes

figure reprinted from www.kdnuggets.com with explicit permission of the editor


Stopping Criteria



- no (“significant”) decrease in $J(W)$

figure reprinted from www.kdnuggets.com with explicit permission of the editor


Using Curvature



- can reduce number of steps
- large curvature: big step

figure reprinted from www.kdnuggets.com with explicit permission of the editor


Using Curvature (cont'd)



- small curvature: small step

figure reprinted from www.kdnuggets.com with explicit permission of the editor


Curvature Trade-off



- use fewer steps
- but more computations

figure reprinted from www.kdnuggets.com with explicit permission of the editor


Issue: Local Minimum



- may consider randomization

figure reprinted from www.kdnuggets.com with explicit permission of the editor


Issue: Noisy Data



- may stop far from optimal

figure reprinted from www.kdnuggets.com with explicit permission of the editor


Issue: Zero Gradient



- cannot update weights

figure reprinted from www.kdnuggets.com with explicit permission of the editor

Issue: Discontinuous Cost Function



- may stop far from optimal

figure reprinted from www.kdnuggets.com with explicit permission of the editor


The Algorithm

- compute w to minimize $f(x)$
- choose initial weights w_0 and learning rate α
- for iteration i update weights:

$$W_i = W_{i-1} - \alpha \cdot \partial f(W_{i-1})$$

- repeat iterations until changes in w do not (significantly) change $f(\cdot)$

Geometric Interpretation



- iterative optimization
- direction to minimize $f(x)$

Python Code

```
import numpy as np
x = np.linspace(0, 10, 1000)
y = (x - 5)**2
df = lambda x: 2*(x-5)

rate = 0.25; precision = 0.001
next_x = 8; max_iterations = 100

for i in range(max_iterations):
    cur_x = next_x
    next_x = cur_x - rate * df(cur_x)
    step = next_x - cur_x
    print("Iteration:", i+1, "x= ", next_x)
    if abs(step) <= precision:
        break

print("Minimum at", next_x, ', ',
      i+1, ' iterations')
```

Execution Details

- for iteration 1:

$$cur_x = next_x = 8$$

$$df(cur_x) = 2 * (8 - 5) = 6$$

$$\begin{aligned} next_x &= cur_x - rate * df(cur_x) \\ &= 8 - 0.25 * 6 = 6.5 \end{aligned}$$

Iteration: 1 x = 6.5

Iteration: 2 x = 5.75

Iteration: 3 x = 5.375


Iteration: 10 x = 5.0029296875

Iteration: 11 x = 5.00146484375

Iteration: 12 x = 5.000732421875

Minimum at 5.000732421875, 12 iterations


Choosing the Learning Rate



- small α - slow convergence
- large α - possible oscillations

figure reprinted from www.kdnuggets.com with explicit permission of the editor

Effect of Decreasing Rate



- slower rate - more iterations

Effect of Decreasing Rate (cont'd)

- $\text{rate} = 0.25$, $\text{next_x} = 8$

Iteration: 1 $x = 7.4$

Iteration: 2 $x = 6.92$


Iteration: 3 $x = 6.536$

Iteration: 29 $x = 5.0046422751473205$

Iteration: 30 $x = 5.003713820117857$

Minimum at 5.003713820117857, 30 iterations

Effect of Increasing Rate



- may fail to converge

Effect of Increasing Rate (cont'd)


- `rate = 1.1`, `next_x = 8`
- may fail to converge

```
Iteration: 1    x = 1.3999999999999995
Iteration: 2    x = 9.32
Iteration: 3    x = -0.184000000000000105
```

```
-----
-----
```

```
Iteration: 99   x = -207044931.30503917
Iteration: 100  x = 248453928.566047
Minimum at 248453928.566047, 100 iterations
```

Oscillations



- never converges

Example of Oscillations

- take $rate = 1, next_x = 6$

- iteration 1:

$$cur_x = next_x = 6$$

$$df(cur_x) = 2 * (6 - 5) = 2$$

$$\begin{aligned} next_x &= cur_x - rate * df(cur_x) \\ &= 6 - 1 * 2 = 4 \end{aligned}$$

- iteration 2:

$$cur_x = next_x = 4$$

$$df(cur_x) = 2 * (4 - 5) = -2$$

$$\begin{aligned} next_x &= cur_x - rate * df(cur_x) \\ &= 4 - 1 * (-2) = 6 \end{aligned}$$

Notes on Gradient Descent

- slow - each iteration requires to examine all samples
- some variants - stochastic gradient descent (examine some points)
- how to compute rate?
 1. constant
 2. decrease with updates

Concepts Check:

- (a) optimization by iterations
- (b) gradient
- (c) curvature
- (d) stopping criteria
- (e) learning rate
- (f) oscillations